

Small and Sweet MapReduce Algorithms

Yufei Tao

Chinese University of Hong Kong

MapReduce

A platform for massive parallel computation.

- First (formal) paper in 2004.
- Now a large collection of algorithms on this platform.

This talk:

Principles in the design and analysis of algorithms in MapReduce
(and other similar platforms).

This talk:

Principles in the design and analysis of algorithms in MapReduce (and other similar platforms).

Ultimate goals:

- **Small:** Simple enough for practical implementation
- **Sweet:** With non-trivial theoretical guarantees.

Computation Model

This is the starting point before any meaningful analysis.

1. MRC model [SODA'10]
2. BSP-like model [ISSAC'11]

3. Minimality Model [SIGMOD'13]
4. Massively Parallel Computation (MPC) [PODS'13].

Computation Model

This talk:

3. Minimality Model [SIGMOD'13]

- A very stringent model – suitable for studying “easy” problems.

4. Massively Parallel Computation (MPC) [PODS'13].

- A more relaxed model – suitable for studying “hard” problems.

Minimality Model

Input and Machines

p = number of machines.

n = number of elements in the input (i.e., dataset)

Assumption: $p \leq \sqrt{\frac{n}{2 \ln n}}$

At the beginning, $O(n/p)$ elements per machine.

- The algorithm has no control over how the elements are initially distributed.

Minimality Model

Superstep

Two phases:

- 1 Machines message each other.

Minimality Model

Superstep

Two phases:

- 1 Machines message each other.
- 2 Machines perform local computation (all operations in the RAM model allowed).

Minimality Model

Superstep

Two phases:

- 1 Machines message each other.
- 2 Machines perform local computation (all operations in the RAM model allowed).

Algorithm

An **algorithm** is a sequence of supersteps.

Minimality Model

Consider a computational problem Π . Suppose that it can be solved in $f(n) = \Omega(n)$ time **sequentially** in the RAM model.

Minimality Model

Consider a computational problem Π . Suppose that it can be solved in $f(n) = \Omega(n)$ time **sequentially** in the RAM model.

Π is **minimal** if it admits an algorithm satisfying all:

- 1 It has $O(1)$ supersteps.
- 2 It uses $O(n/p)$ space at all times.
- 3 Every machine sends $O(n/p)$ words in total.
- 4 Every machine incurs $O(f(n/p))$ CPU time in total.

The algorithm is called a **minimal algorithm** for Π .

Minimality Model

Requirement 1: $O(1)$ supersteps.

Minimality Model

Requirement 1: $O(1)$ supersteps.

Requirement 2: $O(n/p)$ space at all times.

Balanced during the entire algorithm.

Minimality Model

Requirement 1: $O(1)$ supersteps.

Requirement 2: $O(n/p)$ space at all times.

Balanced during the entire algorithm.

Requirement 3: $O(n/p)$ words sent and received.

Asymptotically the same time for each machine to look at its portion of the input.

Minimality Model

Requirement 4: $O(f(n/p))$ CPU time for every machine.

- The MapReduce algorithm **must improve automatically** whenever a faster sequential algorithm is discovered!
 - Implies that the MapReduce algorithm must utilize a sequential algorithm as a **black box**!

Minimality Model

Requirement 4: $O(f(n/p))$ CPU time for every machine.

- The MapReduce algorithm **must improve automatically** whenever a faster sequential algorithm is discovered!
 - Implies that the MapReduce algorithm must utilize a sequential algorithm as a **black box**!
- The best you can do!
 - If you can achieve $o(f(n/p))$ on MapReduce, you can simulate the same algorithm in the RAM model in $o(p \cdot f(n/p))$ time.

Minimal Algorithm for Sorting

At the beginning: n elements on p machines.

At the end: it is possible to label the machines from 1 to p such that all elements on machine i are smaller than those on machine j , for any $i < j$.

Minimal Algorithm for Sorting

At the beginning: n elements on p machines.

At the end: it is possible to label the machines from 1 to p such that all elements on machine i are smaller than those on machine j , for any $i < j$.

Remark (illustration of Requirement 4)


- Comparison-based sort: $\Theta(n \log n)$ sequentially $\Rightarrow O(\frac{n}{p} \log \frac{n}{p})$ MapReduce .
- Sorting n integers: $O(n \log \log n)$ sequentially $\Rightarrow O(\frac{n}{p} \log \log \frac{n}{p})$ MapReduce.

All by the same algorithm.

Minimal Algorithm for Sorting

Superstep 1


- **Phase 1:** Each machine samples each object with probability $\rho = \frac{p}{n} \ln(nt)$ independently, and sends the samples to all other machines.
- **Phase 2:** Each machine sorts the samples received, and identify the **split elements** that partition the list into t segments evenly.
 - Machine i will get all the elements in the i -th segment at the end.



Minimal Algorithm for Sorting

Superstep 2

- **Phase 1:** Each machine sends the elements in the i -th segment to machine i .
- **Phase 2:** Machine i sorts the elements in the i -th segment.



Minimal Algorithm for Sorting

Theorem: The above algorithm is minimal with probability at least $1 - O(1/n)$.

Minimality \Rightarrow MPC

For certain problems, if $O(1)$ supersteps are required, it is impossible to ensure that each machine sends $O(n/p)$ words only.

\Rightarrow **No minimal algorithms.**

Minimality \Rightarrow MPC

For certain problems, if $O(1)$ supersteps are required, it is impossible to ensure that each machine sends $O(n/p)$ words only.

\Rightarrow **No minimal algorithms.**

Phrased differently, if $O(1)$ supersteps are required, we must allow a machine to send more than $O(n/p)$ words – but how much more?

\Rightarrow **The MPC model.**

MPC Model

Input and Machines

p = number of machines.

n = number of elements in the input (i.e., dataset)

Assumption: $p \leq n^{1/c}$ for some constant $c > 1$.

At the beginning, $O(n/p)$ elements per machine.

- The algorithm has no control over how the elements are initially distributed.

MPC Model

Superstep

Two phases:

- 1 Machines message each other.
- 2 Machines perform local computation (all operations in the RAM model allowed).

Algorithm

An **algorithm** is a sequence of supersteps.

MPC Model

An algorithm is required to finish in $O(1)$ supersteps.

MPC Model

An algorithm is required to finish in $O(1)$ supersteps.

Communication cost of one machine =

Total number of words it communicated (sent and received combined)

MPC Model

An algorithm is required to finish in $O(1)$ supersteps.

Communication cost of one machine =

Total number of words it communicated (sent and received combined)

Load of the algorithm =

Communication cost of the most expensive machine

MPC Model

An algorithm is required to finish in $O(1)$ supersteps.

Communication cost of one machine =

Total number of words it communicated (sent and received combined)

Load of the algorithm =

Communication cost of the most expensive machine

Aim: Minimize load.

MPC Model

An algorithm is required to finish in $O(1)$ supersteps.

Communication cost of one machine =

Total number of words it communicated (sent and received combined)

Load of the algorithm =

Communication cost of the most expensive machine

Aim: Minimize load.

Remark 1: CPU time for free.

MPC Model

An algorithm is required to finish in $O(1)$ supersteps.

Communication cost of one machine =

Total number of words it communicated (sent and received combined)

Load of the algorithm =

Communication cost of the most expensive machine

Aim: Minimize load.

Remark 1: CPU time for free.

Remark 2: Every minimal algorithm must have load $O(n/p)$.

MPC Model

An algorithm is required to finish in $O(1)$ supersteps.

Communication cost of one machine =

Total number of words it communicated (sent and received combined)

Load of the algorithm =

Communication cost of the most expensive machine

Aim: Minimize load.

Remark 1: CPU time for free.

Remark 2: Every minimal algorithm must have load $O(n/p)$.

Remark 3: Any problem admits an algorithm of load $O(n)$.

MPC Model

The MPC model has largely been used to study problems where the output size can be much larger than n , e.g., joins.

MPC Model

The MPC model has largely been used to study problems where the output size can be much larger than n , e.g., joins.

Focus of this talk: **Cartesian Product**

MPC Model

The MPC model has largely been used to study problems where the output size can be much larger than n , e.g., joins.

Focus of this talk: **Cartesian Product**

R : a set of n red elements

S : a set of n blue elements

At the beginning: The $2n$ elements are stored on the p machines (each $O(n/p)$ elements).

At the end: Each element in $R \times S$ must appear on some machine.

MPC Lower Bound for Cartesian Product

To warm up, we will prove that, when $p = 2$, any algorithm solving the problem must have load $\Omega(n)$.

⇒ **Implication:** No better algorithm than simply sending everything over to the other machine.

MPC Lower Bound for Cartesian Product

For simplicity, assume each machine has n elements at the beginning.

Suppose that an algorithm solves the problem with load L .

MPC Lower Bound for Cartesian Product

For simplicity, assume each machine has n elements at the beginning.

Suppose that an algorithm solves the problem with load L .

\Rightarrow Machine 1 sees $n + L$ elements overall.

MPC Lower Bound for Cartesian Product

For simplicity, assume each machine has n elements at the beginning.

Suppose that an algorithm solves the problem with load L .

⇒ Machine 1 sees $n + L$ elements overall.

⇒ It can produce at most

$$\left(\frac{n + L}{2}\right)^2$$

pairs

MPC Lower Bound for Cartesian Product

For simplicity, assume each machine has n elements at the beginning.

Suppose that an algorithm solves the problem with load L .

⇒ Machine 1 sees $n + L$ elements overall.

⇒ It can produce at most

$$\left(\frac{n + L}{2} \right)^2$$

pairs

⇒ Same for machine 2.

MPC Lower Bound for Cartesian Product

For simplicity, assume each machine has n elements at the beginning.

Suppose that an algorithm solves the problem with load L .

⇒ Machine 1 sees $n + L$ elements overall.

⇒ It can produce at most

$$\left(\frac{n + L}{2}\right)^2$$

pairs

⇒ Same for machine 2.

⇒

$$2 \left(\frac{n + L}{2}\right)^2 \geq n^2$$

MPC Lower Bound for Cartesian Product

For simplicity, assume each machine has n elements at the beginning.

Suppose that an algorithm solves the problem with load L .

⇒ Machine 1 sees $n + L$ elements overall.

⇒ It can produce at most

$$\left(\frac{n + L}{2}\right)^2$$

pairs

⇒ Same for machine 2.

⇒

$$2 \left(\frac{n + L}{2}\right)^2 \geq n^2$$

which solves to $L \geq \frac{2-\sqrt{2}}{2\sqrt{2}} n$.

MPC Lower Bound for Cartesian Product

Next, we will extend the algorithm to prove that, any cartesian product algorithm must have load $\Omega(n/\sqrt{p})$.

⇒ **Implication:** No minimal algorithm!

MPC Lower Bound for Cartesian Product

Suppose that an algorithm has load L .

MPC Lower Bound for Cartesian Product

Suppose that an algorithm has load L .

\Rightarrow Every machine can receive at most L elements.

MPC Lower Bound for Cartesian Product

Suppose that an algorithm has load L .

⇒ Every machine can receive at most L elements.

⇒ In total, the machine sees $c \cdot n/p + L$ elements during the entire algorithm.

MPC Lower Bound for Cartesian Product

Suppose that an algorithm has load L .

⇒ Every machine can receive at most L elements.

⇒ In total, the machine sees $c \cdot n/p + L$ elements during the entire algorithm.

⇒ The machine can produce at most

$$\left(\frac{c \cdot n/p + L}{2} \right)^2$$

pairs.

MPC Lower Bound for Cartesian Product

Suppose that an algorithm has load L .

⇒ Every machine can receive at most L elements.

⇒ In total, the machine sees $c \cdot n/p + L$ elements during the entire algorithm.

⇒ The machine can produce at most

$$\left(\frac{c \cdot n/p + L}{2} \right)^2$$

pairs.

⇒ Putting together p machines means

$$p \cdot \left(\frac{c \cdot n/p + L}{2} \right)^2 \geq n^2$$

MPC Lower Bound for Cartesian Product

Suppose that an algorithm has load L .

⇒ Every machine can receive at most L elements.

⇒ In total, the machine sees $c \cdot n/p + L$ elements during the entire algorithm.

⇒ The machine can produce at most

$$\left(\frac{c \cdot n/p + L}{2} \right)^2$$

pairs.

⇒ Putting together p machines means

$$p \cdot \left(\frac{c \cdot n/p + L}{2} \right)^2 \geq n^2$$

which solves to $L = \Omega(n/\sqrt{p})$.

MPC Lower Bound for Pairwise Natural Join

Pairwise Natural Join

$R = (A, B)$: a table of n tuples

$S = (B, C)$: a table of n tuples

At the beginning: The $2n$ tuples are stored on the p machines.

At the end: Each element in $R \bowtie S$ must appear on some machine.

MPC Lower Bound for Pairwise Natural Join

Pairwise Natural Join

$R = (A, B)$: a table of n tuples

$S = (B, C)$: a table of n tuples


At the beginning: The $2n$ tuples are stored on the p machines.

At the end: Each element in $R \bowtie S$ must appear on some machine.

LOWER BOUND: Load $\Omega(n/\sqrt{p})!$

MPC Cartesian Product

An optimal algorithm—named **cube**—with load $O(n/\sqrt{p})$ is known. The algorithm actually computes the cartesian product.



MPC Model: Understood Problems

Problems whose communication complexities (i.e., load) have been well understood:

- Enumeration of constant-sized subgraphs.
- Restricted natural joins.
- Linear programming.
- ...

Summary

Principles in design and analysis of MapReduce algorithms

- Minimality
- MPC

Summary

Principles in design and analysis of MapReduce algorithms

- Minimality
- MPC

Future work 1: Re-examine **all** problems!

Summary

Principles in design and analysis of MapReduce algorithms

- Minimality
- MPC

Future work 1: Re-examine **all** problems!

Future work 2: Lower bound for problems with **small** result sizes!

- From the theory of communication complexity, it is easy to prove that when $p = 2$, $\Omega(n)$ communication is necessary for detecting whether a natural join has an empty result!

Summary

Principles in design and analysis of MapReduce algorithms

- Minimality
- MPC

Future work 1: Re-examine **all** problems!

Future work 2: Lower bound for problems with **small** result sizes!

- From the theory of communication complexity, it is easy to prove that when $p = 2$, $\Omega(n)$ communication is necessary for detecting whether a natural join has an empty result!

Future work 3: $o(n/p)$ -load algorithms!