

# Introduction to Computation Models

Yufei Tao

Department of Computer Science and Engineering  
Chinese University of Hong Kong

Computer science is a subject of mathematics.


The subject studies how to apply **atomic operations** to accomplish a task with the lowest **cost**.

Different **computation models** produce different branches of the subject.

Today, we will look at several mainstream models.

## I. Random Access Machine (RAM)

This is the model you used in your undergraduate study.



**Atomic operations** ( $a$  and  $b$  are registers):

- $a \leftarrow 100, a \leftarrow b$
- $a + b, a - b, a * b, a / b$
- $a < b, a = b, \text{ or } a > b$
- read a memory cell into  $a$ , or write  $a$  into a memory cell.

**Time** of an algorithm = # atomic operations performed.

**The Sorting Problem:** Given an array  $A$  of  $n$  distinct integers, produce another array where the same integers have been arranged in ascending order.

## Merge Sort

- **Divide:** Let  $A_1$  be the array containing the first  $\lceil n/2 \rceil$  elements of  $A$ , and  $A_2$  be the array containing the other elements of  $A$ . Sort  $A_1$  and  $A_2$  recursively.
- **Conquer:** Merge the two sorted arrays  $A_1$  and  $A_2$  in ascending order. This can be done in  $O(n)$  time.

$f(n)$ : time of sorting  $n$  integers

$$f(n) = 2 \cdot f(n/2) + O(n)$$

which yields  $f(n) = O(n \log n)$ .

A computation model falls short when it can no longer be used to guide the design of algorithms.


Today, a modern CPU has multiple cores and supports **parallel execution**.

RAM can no longer capture an algorithm's behavior on those CPUs.

## II. Parallel Random Access Machine (PRAM)



$p = \# \text{ CPUs}$




A **step**: each CPU **independently** performs an atomic operation ( $a$  and  $b$  are registers inside the CPU):

- $a \leftarrow 100, a \leftarrow b$
- $a + b, a - b, a * b, a / b$
- $a < b, a = b, \text{ or } a > b$
- read a memory cell into  $a$ , or write  $a$  into a memory cell.

**Constraint:** In a step, if a CPU writes to a cell, no other CPU can read or write to the cell.

This is the **CREW** model (common read exclusive write).

**The Sorting Problem:** Given an array  $A$  of  $n \geq p$  distinct integers, produce another array where the same integers have been arranged in ascending order.



Naive:  $O(n \log n)$  steps.

Optimal:  $O(\frac{n}{p} \log n)$  steps. **Why?**

We will see how to do  $O(\frac{n}{p} \log^2 n)$  steps.

We will assume  $p = n$ . Extending the discussion to  $p < n$  is left to you.

**The Merging Problem:** Let  $A_1$  and  $A_2$  be two sorted arrays, each with  $p/2$  integers. Merge  $A_1$  and  $A_2$  into one sorted array (with  $p$  integers).

Naive:  $O(p)$  steps.

We will do  $O(\log p)$  steps.

$A_1$ 

17	26	28	38	41	72	83	88
----	----	----	----	----	----	----	----

 $A_2$ 

5	9	12	35	47	52	68	69
---	---	----	----	----	----	----	----

Consider the number 26 in  $A_1$ .

- It is the 2nd smallest in  $A_1$ .
- It is greater than 3 numbers in  $A_2$ .

Hence, 26 should be the  $2 + 3 = 5$ -th in the merged array.

$A_1$							
17	26	28	38	41	72	83	88
$B_1$							
4	5	6	8	9	14	15	16

$A_2$							
5	9	12	35	47	52	68	69
$B_2$							
1	2	3	7	10	11	12	13

We want to generate arrays  $B_1$  and  $B_2$ . For each number  $x$  in  $A_1$  (or  $A_2$ ), the corresponding cell in  $B_1$  (or  $B_2$ ) indicates the position of  $x$  in the final merged array.

Once  $B_1$  and  $B_2$  are ready, the merged array can be produced in  $O(1)$  steps — recall that we have  $p = 16$  CPUs.

$$A_1$$

17	26	28	38	41	72	83	88
----	----	----	----	----	----	----	----

	5						
--	---	--	--	--	--	--	--

$$B_1$$

$$A_2$$

5	9	12	35	47	52	68	69
---	---	----	----	----	----	----	----

How to generate  $B_1$  and  $B_2$ ?

By assigning one CPU to each number, we can generate the number's position indicator in  $O(\log p)$  steps using binary search.

This gives an algorithm solving the merging problem in  $O(\log p)$  steps.

We can now sort  $n$  elements with  $n$  CPUs in  $O(\log n)$  phases.

- Phase 0:  $n/2$  parallel merges, each merging two lists of size 1.
- Phase 1:  $n/4$  parallel merges, each merging two lists of size 2.
- Phase 2:  $n/8$  parallel merges, each merging two lists of size 4.
- ...

Each phase takes  $O(\log n)$  steps.

Overall,  $O(\log^2 n)$  steps.

### III. External Memory (EM)




RAM and PRAM assume that the data fit in memory.

In practice, a dataset often needs to be stored in the disk. In those scenarios, CPU operations are rarely the performance bottleneck. Typically, an algorithm's efficiency depends on how many **disk I/Os** are performed.

This motivates the **external memory** model.

$M$  = the number of words in memory

$B$  = the number of words in a disk block



## A I/O

- reads a disk block into memory, or
- writes  $B$  memory words into a disk block.

**Cost** of an algorithm: # I/Os performed.

CPU calculation is for free.

**The Sorting Problem:** Given a file  $A$  of  $n \geq B$  distinct integers stored in  $O(n/B)$  blocks, produce another file of  $O(n/B)$  blocks where the same integers have been arranged in ascending order.

Naive:  $O(n \log n)$  I/Os.


We will do:  $O(\frac{n}{B} \log_{M/B} \frac{n}{B})$  I/Os.

**The Merging Problem:** Let  $A_1, A_2, \dots, A_{\frac{M}{B}-1}$  be  $\frac{M}{B} - 1$  files such that


- each file stores  $t_i \geq B$  integers in ascending order with  $O(t_i/B)$  block;
- the integers in all the files are distinct.

Let  $t = \sum_{i=1}^{M/B-1} t_i$ . Produce another file of  $O(t/B)$  blocks where all the  $t$  integers have been arranged in ascending order.



We will see how to solve the problem in  $O(t/B)$  I/Os.





1 5	2 9	4 6	10 17




1 5	2 9	4 6	10 17










1 5	2 9	4 6	10 17

1 2
4 5



The total cost is  $O(t/B)$  because

- there are  $O(t/B)$  blocks in the input files;
- there are  $O(t/B)$  blocks in the output file.

Think: equipped with the above merging algorithm, how can we solve the sorting problem in  $O(\frac{n}{B} \log_{M/B} \frac{n}{B})$  I/Os?

## IV. Massively Parallel Computation (MPC)

Today, many computational tasks are performed on datasets whose sizes are at the tera-byte order or higher. Massively-parallel systems (such as Spark) are often deployed to carry out those tasks. Such a system involves dozens or hundreds of machines interconnected by a broadband network. CPU time and I/O cost are no longer the performance bottleneck. Instead, the efficiency of an algorithm is determined by **network communication**.

This motivates the **massively parallel computation** (MPC) model.

$p = \#$  machines

A **superstep** includes two steps.

- Step 1: each machines performs local computation.
- Step 2: each machine sends a message to every other machine.

**Cost of a superstep:** the maximum number of words communicated (i.e., sent and received) by a machine in Step 2.

- Step 1 is for free.

**Cost of an algorithm:** total cost of all the supersteps.

### Example:

$p = 3$  machines. Suppose that in a superstep

- machine 1 sends  $50m$  and  $30m$  words to machine 2 and 3, respectively;
- machine 2 sends  $0m$  and  $40m$  words to machine 1 and 3, respectively;
- machine 3 sends  $10m$  and  $5m$  words to machine 1 and 2, respectively.

Cost of the superstep =  $85m$  words.

**The Sorting Problem:** Let  $S$  be a set of  $n$  numbers such that each machine stores  $O(n/p)$  numbers originally. Re-arrange the numbers such that

- each machine still stores  $O(n/p)$  numbers;
- for any  $1 \leq i < j \leq p$ , all the numbers on machine  $i$  are less than those on machine  $j$ .

We will consider  $n \geq p^3$  (a reasonable assumption in practice).

Naive: cost  $O(n)$ .

We will do: cost  $O(n/p)$  in two supersteps.



$S_i$  = the set of numbers on machine  $i \in [1, p]$

### Superstep 1

- **Step 1:** Machine  $i$  samples each integer in  $S_i$  with probability  $\frac{p}{n} \ln(np)$  independently.
- **Step 2:** Each machine sends the samples to all other machines.

machine 1  
(10, 33, 57, 60, 80)

machine 2  
(5, 22, 48, 56, 90)

machine 3  
(15, 35, 38, 88, 96)

---

(5, 33, 38, 48, 80, 88)

(5, 33, 38, 48, 80, 88)

(5, 33, 38, 48, 80, 88)


Let  $s$  be the total number of samples (from all machines)

## Superstep 2

- **Step 1:** Each machine divides  $\mathbb{R}$  into  $p$  disjoint **segments**, each covering at most  $s/p$  samples. Let these segments be  $\sigma_1, \sigma_2, \dots, \sigma_p$  in ascending order.

(5, 33, 38, 48, 80, 88)    (5, 33, 38, 48, 80, 88)    (5, 33, 38, 48, 80, 88)

- **Step 2:** Machine  $i$  sends  $S_i \cap \sigma_j$  to machine  $j$ , for each  $j \in [1, p]$ .



We can prove that the cost of the algorithm is  $O(n/p)$  with probability at least  $1 - O(1/n)$ .

The proof is omitted from this talk.